

The Squawk Java Virtual Machine

Doug Simon

Sun Labs

Overview

- Context and motivation
- VM architecture & implementation
- ARM port details
- Lessons learnt
- Future work

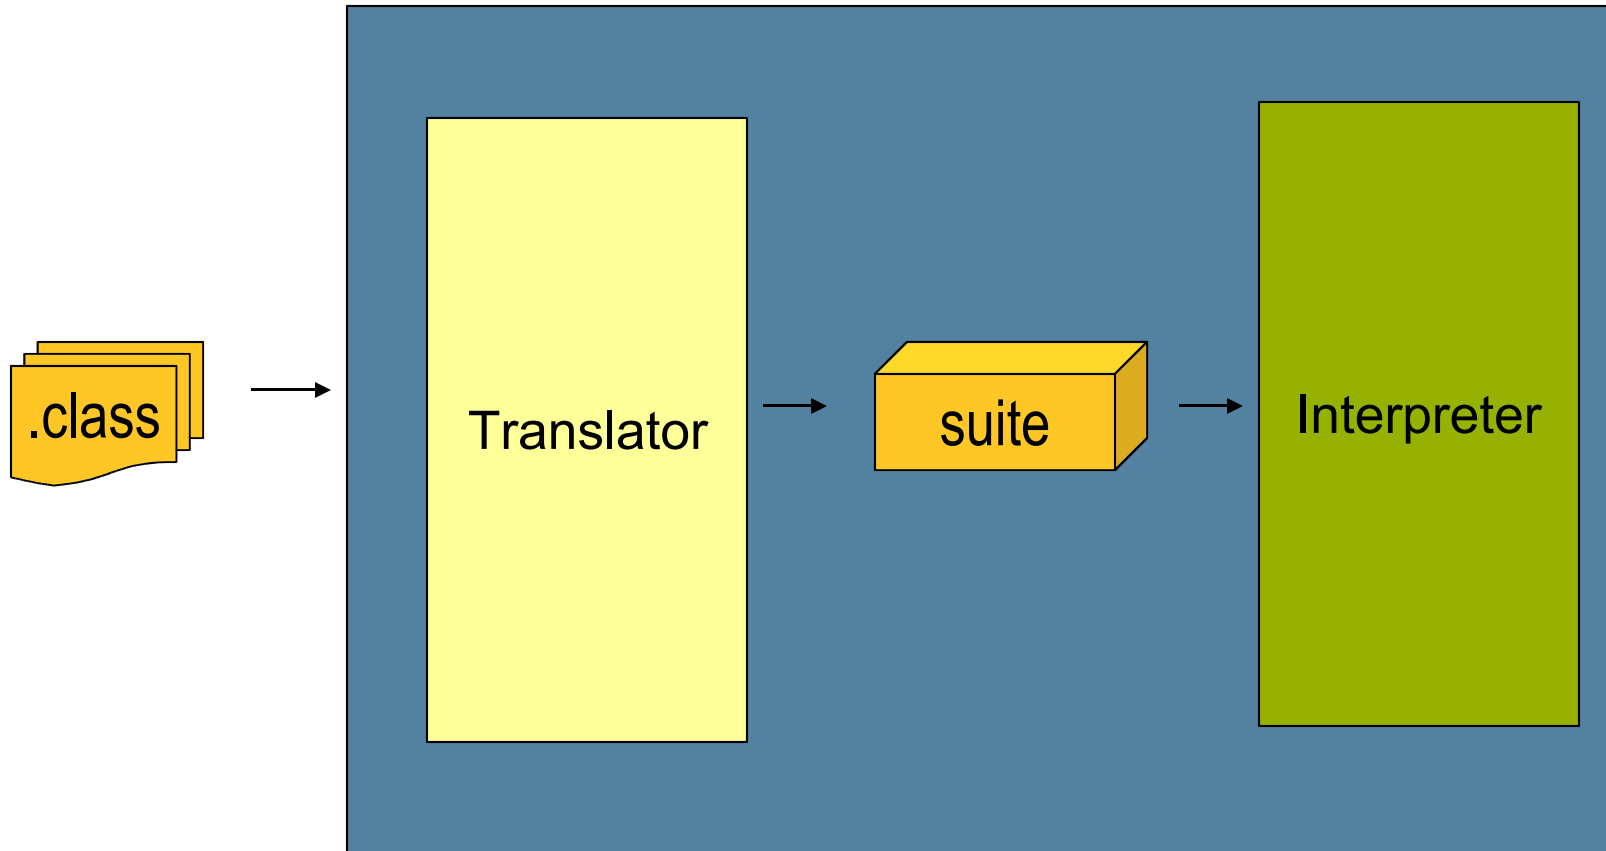
Context and motivation

- Spotless & KVM experience
- Alternative way to construct JVMs?
 - > Do it in Java!
 - > Pointer safety, exceptions, GC...
 - > More portable
 - > Ease of development
 - > Extend loader/verifier to simplify execution
- What does this approach enable?
- Goals:
 - > CLDC implementation for memory constrained devices
 - > Can run on bare metal (i.e. no OS required)

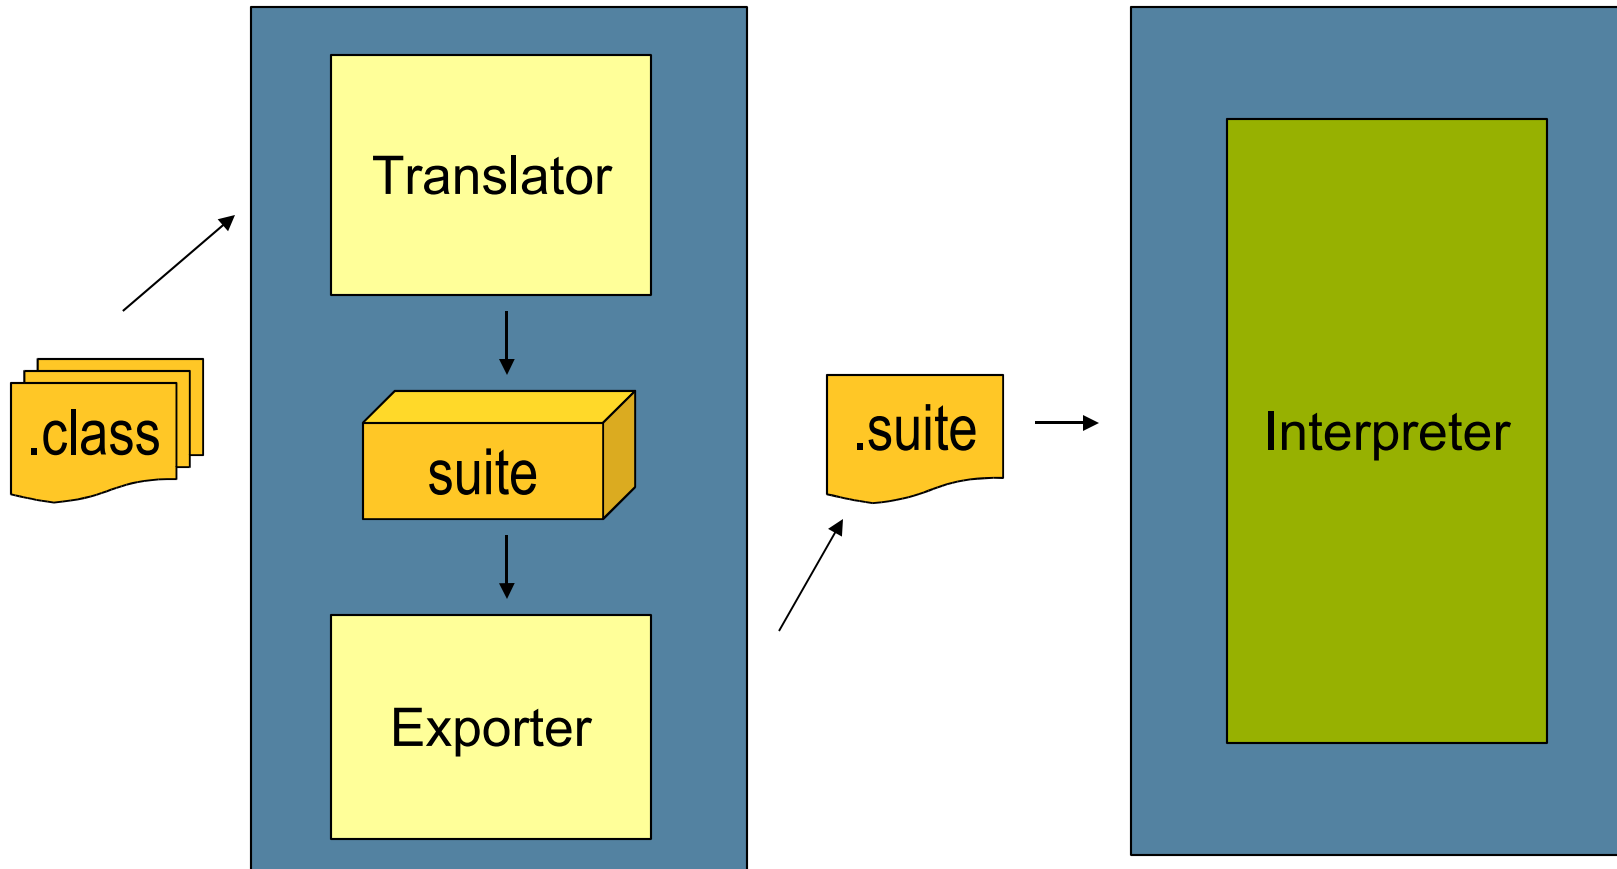
Architecture

- Native
 - > Interpreter
 - > Platform IO
- Java
 - > Translator
 - > Loader, verifier & bytecode transformation
 - > Core
 - > Thread scheduler
 - > Exception handling
 - > Synchronization
 - > Garbage collector
 - > Exporter
 - > JDWP Debug Agent
 - > CLDC API

Architecture



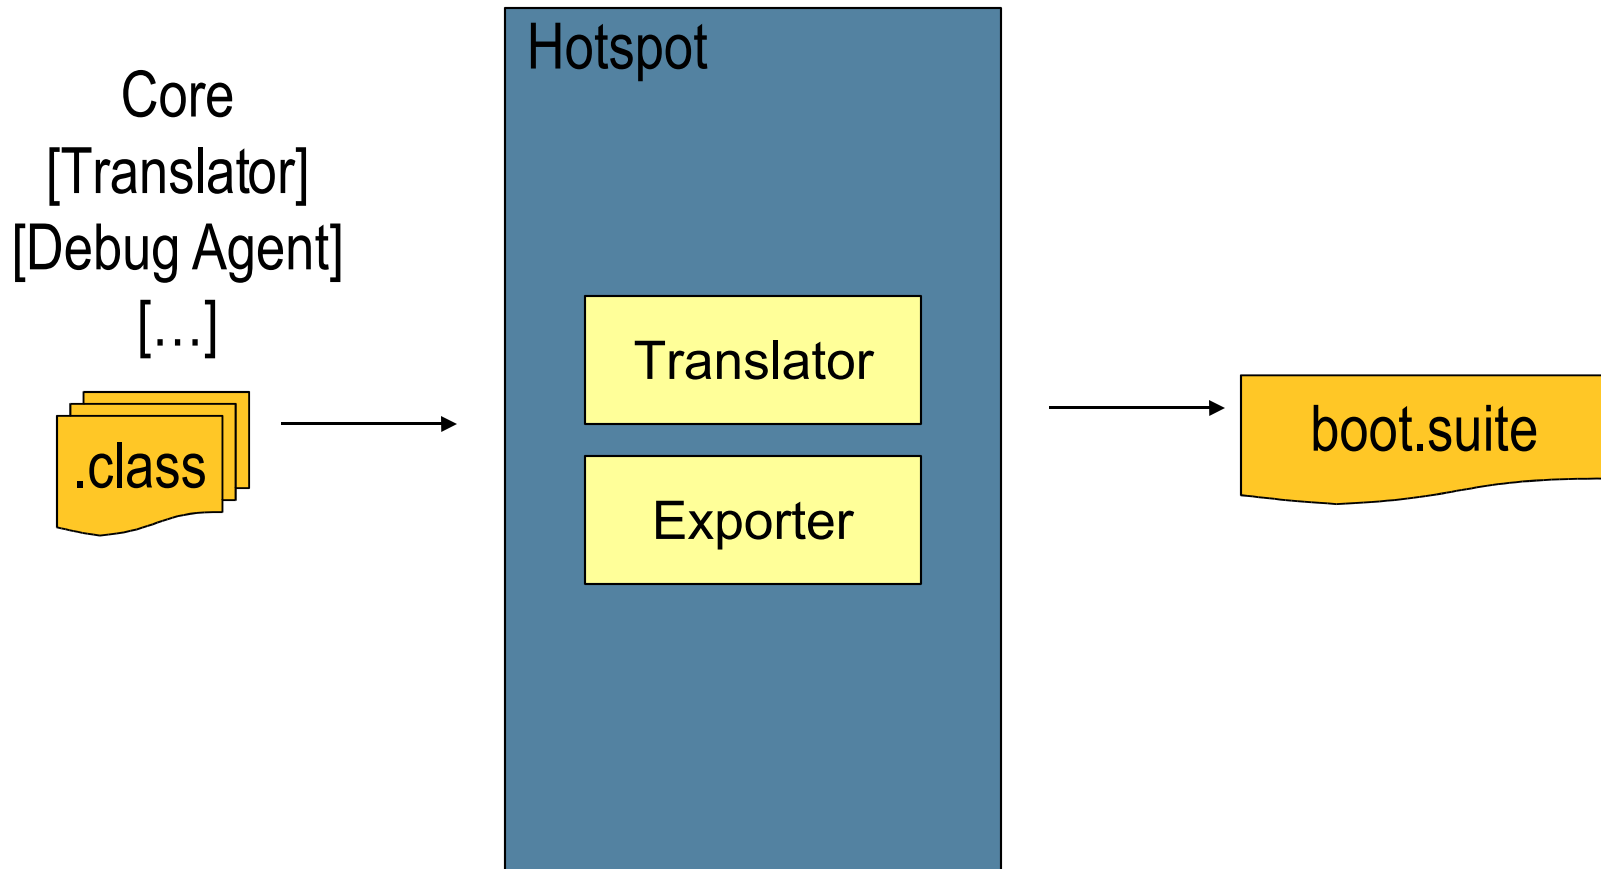
Split VM Architecture



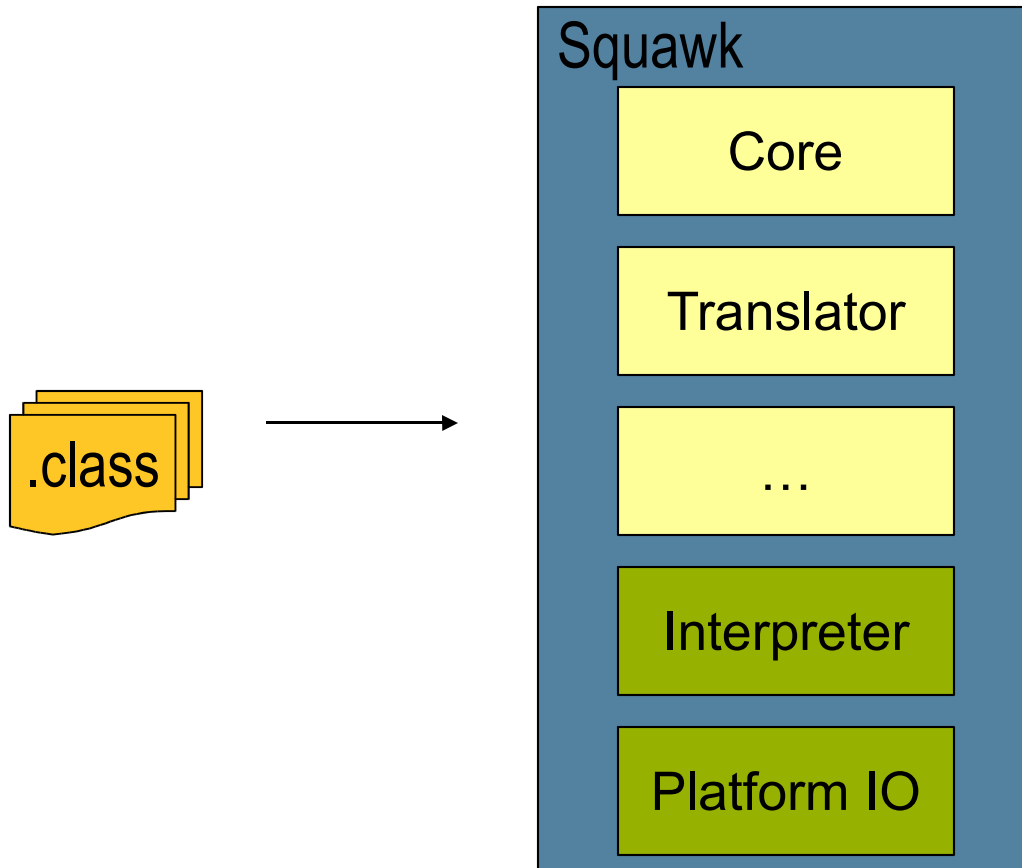
Architecture

- Different configurations for different needs
- Bootstrapping
 - > Used to create boot .suite file
- Standard
 - > CLDC compliant JVM
 - > Used to create (non-boot) .suite files
- Minimal
 - > Minimizes static and dynamic on-device memory requirements

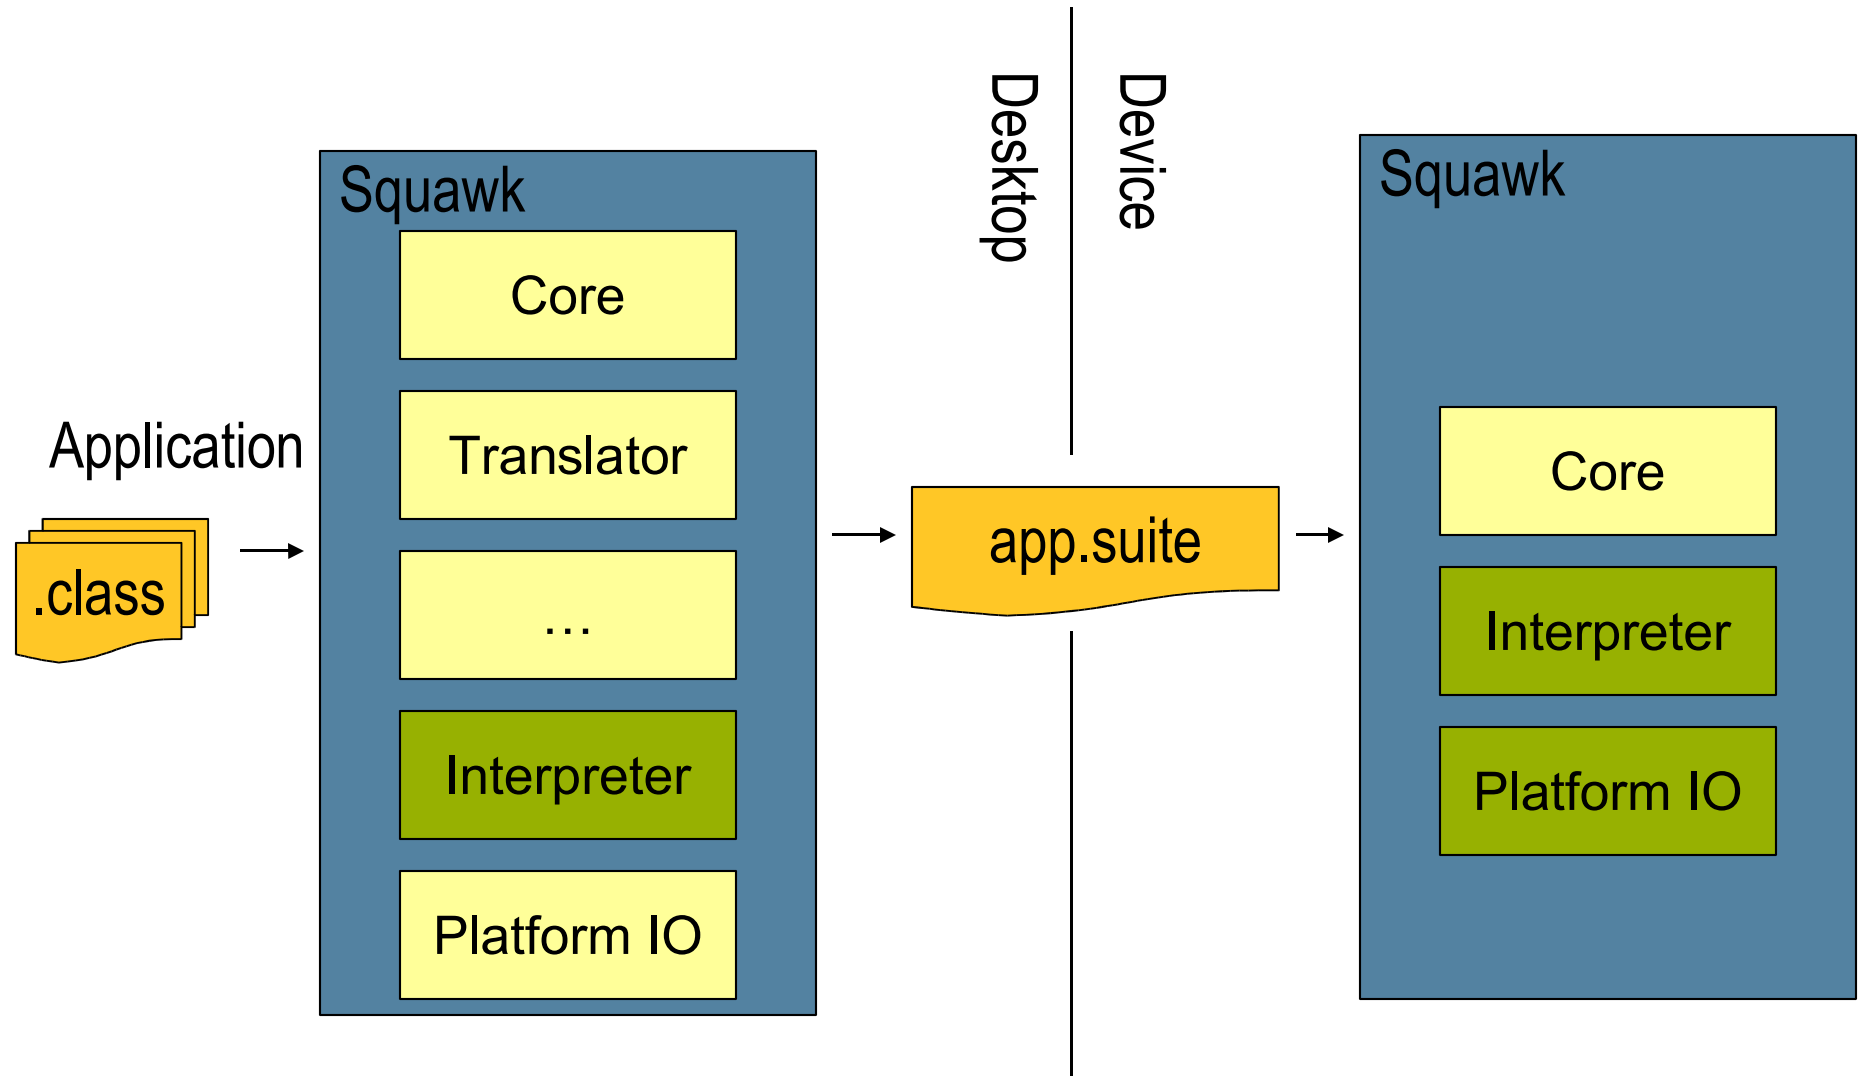
Bootstrapping



Standard




Minimal



Suites

- A suites is a preprocessed set of class files
- Internally fully linked
 - > Pointers to other classes in suite or parent(s) only
 - > Chain of suites is a transitive class closure
- Internal bytecode set optimized for:
 - > Space
 - > Simplified garbage collection
 - > Compilation
 - > In-place execution
 - > Requires pre-linking
 - > Enables deployment in ROM
 - > Improves startup time

Suite chain



Space optimizations


- 2 byte branch
- More 1 byte `load/store/const` instructions
- 2 byte field access
 - > Operand is object offset, not constant pool index
- 2 byte invoke
 - > Operand is vtable index, not constant pool index
- Escape mechanism for:
 - > Float and double instructions
 - > Large operands (i.e. widened operands)

Uncompressed JAR vs Suite File Size Comparison

| Application | JAR | Suite | Suite/JAR |
|---------------|----------------|----------------|-------------|
| CLDC | 458,291 | 149,542 | 0.33 |
| cubes | 38,904 | 16,687 | 0.42 |
| hanoi | 1,805 | 835 | 0.46 |
| delta blue | 30,623 | 8,144 | 0.27 |
| mpeg | 100,917 | 54,888 | 0.54 |
| manyballs | 12,017 | 6,100 | 0.51 |
| pong | 17,993 | 7,567 | 0.42 |
| spaceinvaders | 50,854 | 25,953 | 0.51 |
| tilepuzzle | 18,516 | 7,438 | 0.40 |
| wormgame | 23,985 | 9,131 | 0.38 |
| Total | 753,905 | 286,285 | 0.38 |

Simplified garbage collection

- One pointer map per method
 - > Obviates need for stack maps and analysis during collection
 - > Requires extra slots
 - > Can be mitigated by slot re-allocation based on liveness analysis (instead of lexical scoping)




Activation records


Simplified garbage collection

- Empty operand stack at GC points
 - > Same pros and cons

Normal Java



Squawk




Object memory serialization

- A graph of objects can be serialized to a connection (e.g. file, socket, ...)
- Implemented as an extension to the garbage collector
- Pointers in serialized suites are normalized
- Can swap endianness on load and/or save
- Used by exporter and...

Object layout

- 1 word objects headers, 2 words for arrays



- Interposed ObjectAssociation for hashcode/monitor
 - > Created on demand
- Associative monitor used instead for ROM objects
 - > Rarely needed
 - > Hashcode is address

Method objects

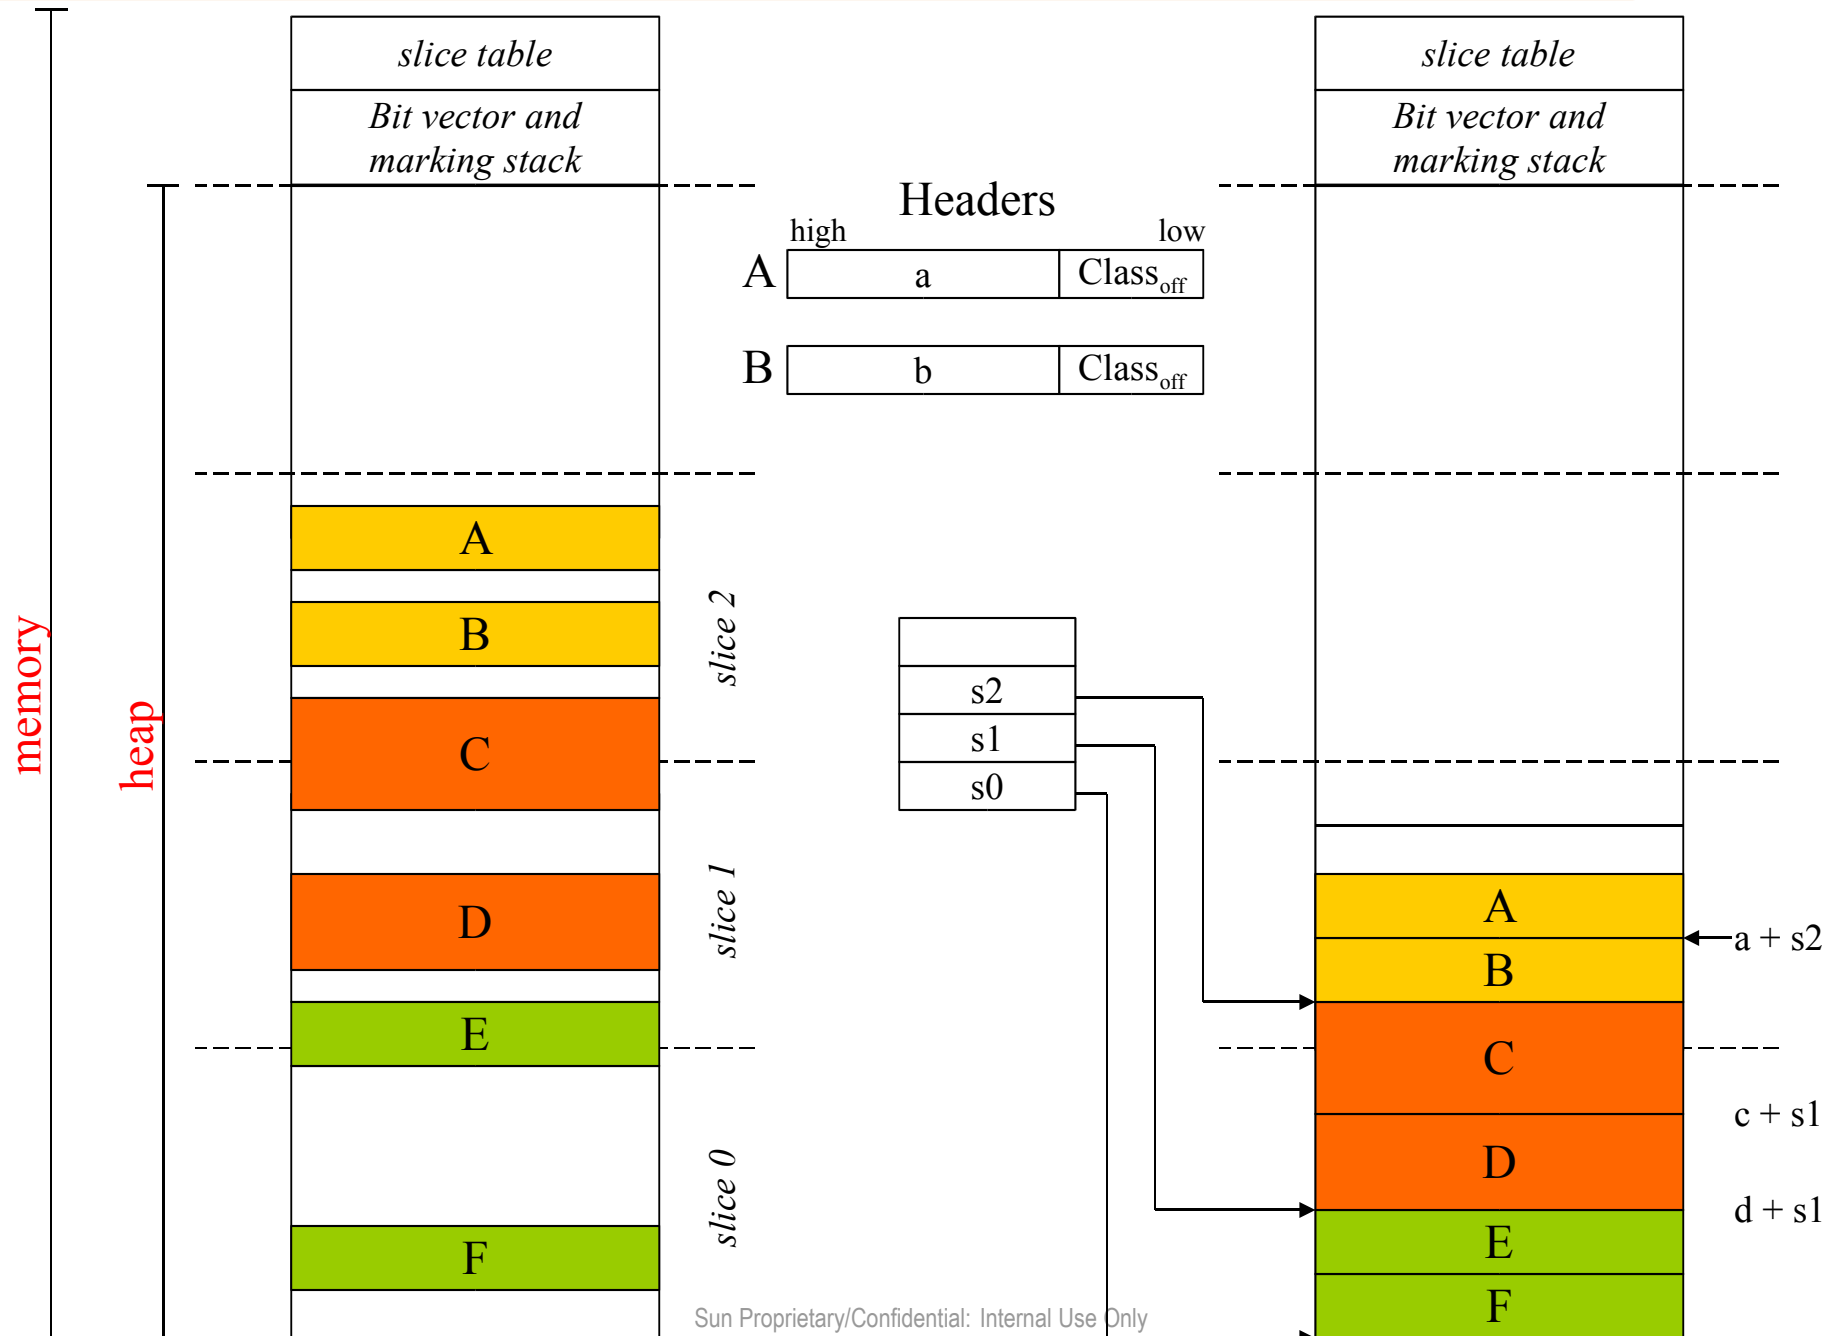
- Specialized byte array with variable length header
- Header encodes:
 - > Defining class pointer
 - > Pointer map
 - > Type map (for non-object and non-int slots)
 - > Used by compiler
 - > Exception handler table
 - > Slots required for parameters, locals, operand stack
- Tighter encoding for common case:
 - > No handlers, parameters<32, locals<32, stack<32
 - > 4 words

Space optimized symbols and metadata

- Symbols for a class encoded in a byte array
 - > Method signatures
 - > Field signatures
 - > Access flags
- Method body metadata stored separately
 - > Line number tables
 - > Local variable tables
- Can strip:
 - > Everything
 - > Private and package private members (library)
 - > Private (extendable library)

Garbage collector

- M&C collector is based on lisp2 algorithm in Lins & Jones
- Generational - sliding window for young gen
- Slices (a la Monty) obviates need for extra header word for forwarding pointer
- Bit vector for entire heap (3% overhead):
 - > Mark bits for collection space
 - > Write-barrier bits for old generation
- Stop-the-world!
- Can copy an object graph into a contiguous block - object serialization




Isolates

- Abstraction for an application in a JVM
- Multiple isolates supported in single Squawk VM
- Objects are:
 - > Logically isolated for each isolate
 - > Physically interleaved on heap
- Non-shared class state:
 - > Static fields
 - > Class initialization state
 - > Class monitors
- Immutable state is shared:
 - > Methods, string constants, classes (apart from above)


JVM Isolates and OS Processes Analogy

Operating System

Process




Process




Squawk JVM

Isolate




Isolate



Shared memory for multiple isolates


Standard JVM



Squawk JVM

Non-shareable object memory

Shareable object memory



Squawk Isolates

- Similar to JSR 121
 - > No support for inter-isolate communication
- Can be paused and resumed
- Can be migrated
 - > Uses same object serialization mechanism as the Exporter
 - > Constraints on external state
 - > Must be none, OR
 - > Must be homogenous at both ends, OR (Sun SPOT)
 - > Must be serializable (desktop Squawk)

Thread scheduling

- Green threads
- Supports blocking in native methods
 - > Thread blocked on event queue polled by scheduler
 - > Events correspond to interrupt
- Rescheduling done at backward branches
 - > `bbtarget` inserted by translator
- System code is non-preemptible
 - > Simplifies VM greatly
 - > Assumes most execution is in application code
- Requires explicit `Thread.yield()` in potentially long running system code
 - > E.g. `System.arraycopy()`

Compiler Interface

- Stack based API with C-like semantics
- Two categories of instructions:
 - > General purpose
 - > `add()`, `sub()`, `eq()`, `push()`, `call()`
 - > Interpreter generator specific
 - > `alloca()`, `getIP()`, `getFP()`
- Intended to be used for:
 - > AOT compiling of boot suite
 - > JIT
 - > Interpreter generator

Prototype Compiler Implementation

- One-pass compiler
 - > Uses shadow-stack representation
- Simple register allocator
- Implemented for x86 and ARM
- TBD: integration with rest of system
- Collaboration with U Sydney to develop optimizing implementation
 - > Probably too heavyweight for JIT but suitable for AOT
- Experimental approach: results pending!

Porting to the ARM

- Squawk ported to the ARM simulator in a week by someone external to VM team
- Within 3 months, 2 contractors had it running on the Atmel ARM board
 - > Including the ARM Platform IO support
- Complete system running on the Sun SPOT within 6 months
 - > Including sensor libraries, demo apps and build notes

Sun SPOT: Small Programmable Object Technology

Sun SPOT: Hardware

- Hardware
 - > 32-bit ARM core
 - > Chipcon CC2420 based wireless platform
 - > 256K RAM, 2Mb FLASH
 - > SPI based peripherals
 - > Simple sensor board
 - > 3D accelerometer
 - > Light sensor
 - > 3 color LEDs
 - > 9 I/O pins
 - > Temperature

Sun SPOT: Software

- Minimal Squawk VM configuration
 - > 350K static footprint:
 - > 80K - Native + PlatformIO
 - > 270K - Core + CLDC 1.0 library (bootstrap suite)
- Device drivers written in Java
- Libraries for
 - > Driving hardware: radio, sensor boards, ...
 - > Basic 802.15.4 network functionality

Sun SPOT + Squawk: Why?

- Improve accessibility of sensor devices
- Motes + nesC was considered hard to program
- Java libraries alleviate a lot of infrastructure programming
- Can use standard Java tools including a debugger!

Lessons & Observations

- Problem: want to write highly factored VM code AND want great performance
- Solution: Need optimizing AOT compiler!
- Annotations useful for expressing VM specific constraints:
 - > method must be inlined (can't invoke virtual)
 - > method must never not do allocation
 - > public class must only be accessed by VM code

Lessons & Observations

- Writing GC in Java was a double-edged sword
 - + No need to duplicate definitions in C
 - Hidden constraints: e.g. can't invoke virtual methods on object with corrupt header
 - Very slow without AOT!
- Don't couple VM code with standard core classes and namespaces
 - > Complicates bootstrapping
 - > Requires -Xbootclasspath and (even worse) -Dsun.boot.library.path (works differently in JDK 5.0!)

Future Work

- Optimize Sun SPOT deployment:
 - > Improve interrupt handling latency
 - > Integrate optimizations in the translator
- Complete the compiler and to validate (or otherwise) the compiler interface approach
- Investigate resource management based on isolates
- Port to a cell phone?
- Open source