

Squawk Java Virtual Machine Compiler Architecture

David Liu

Manager: Cristina Cifuentes

Presentation Outline

- What is Squawk?
- Standard JVM vs Squawk JVM
- Squawk architecture
- Compiler interface
- Multiple compilers - simple vs optimizing
- Execution of Squawk code
- Future work

What is Squawk?

- Small, execute-in-place JVM
- Written in Java, not C
- Optimized for small devices
- Squawk architecture has an interpreter, compiler and jitter
- Current use: Sun Spot sensors platform
- Why write a JVM in Java?
 - > Portability
- What performance can the Squawk JVM achieve?

Standard JVM vs Squawk JVM

Standard JVM

Java class library

Loader

Verifier

Garbage collector

Interpreter

Thread Scheduler

Dynamic Compiler

I/O library

Native code

Java code

C code

Squawk JVM

Java class library

Loader

Verifier

Transformer

Garbage collector

Interpreter

Thread Scheduler

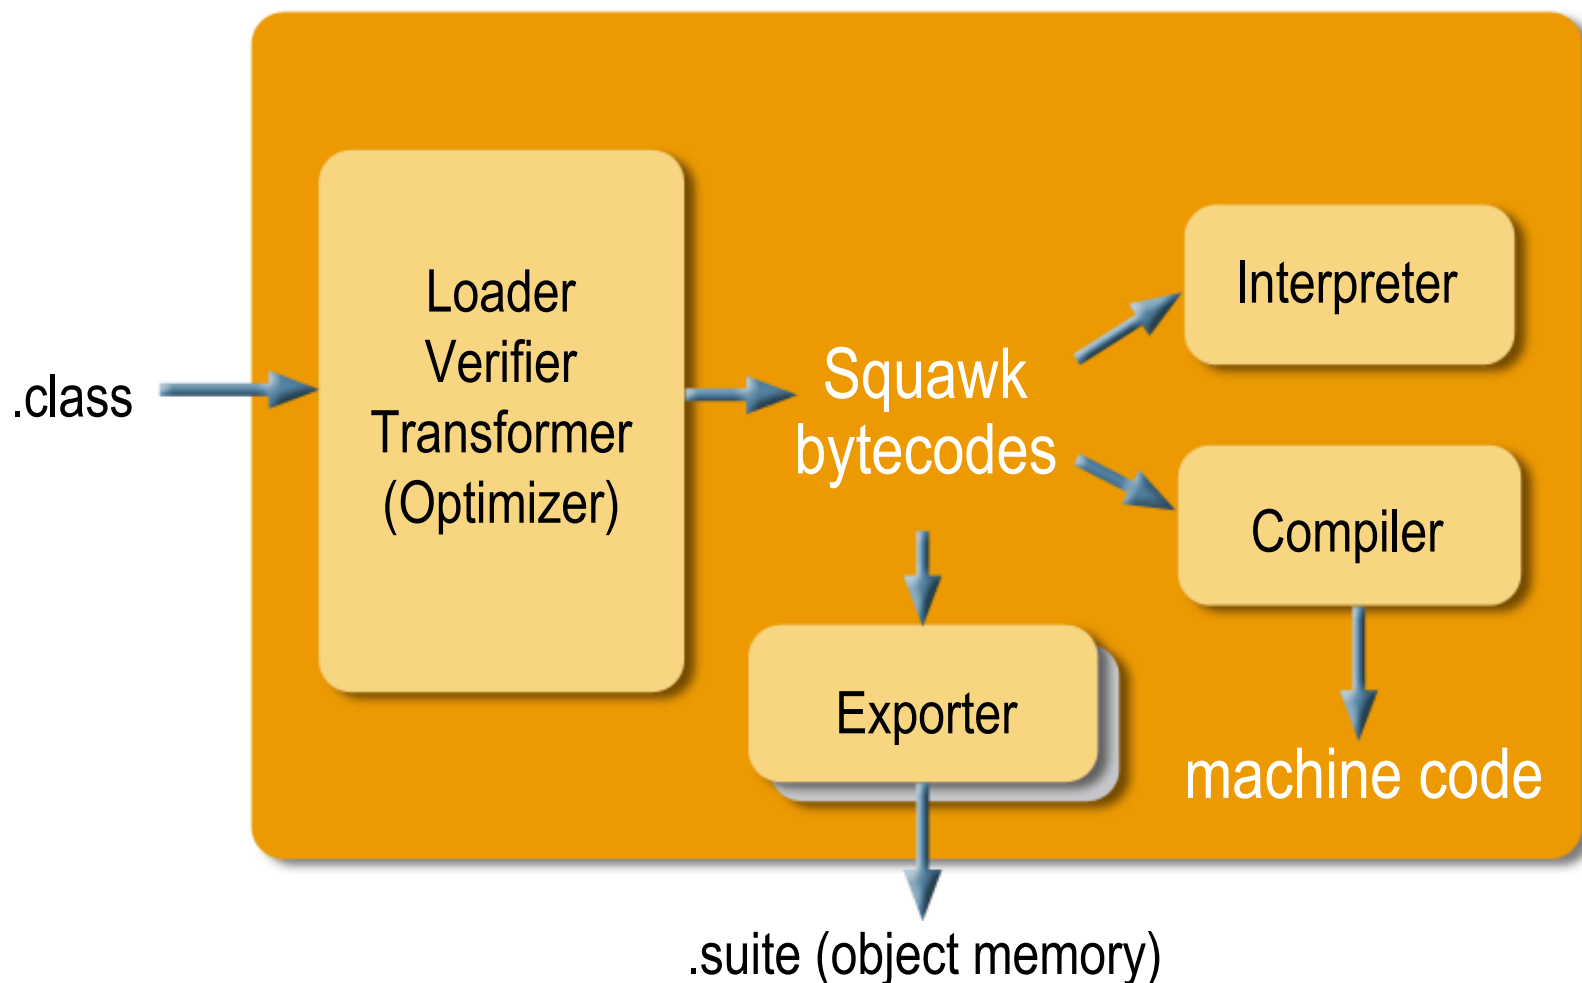
Exporter

Dynamic Compiler

I/O library

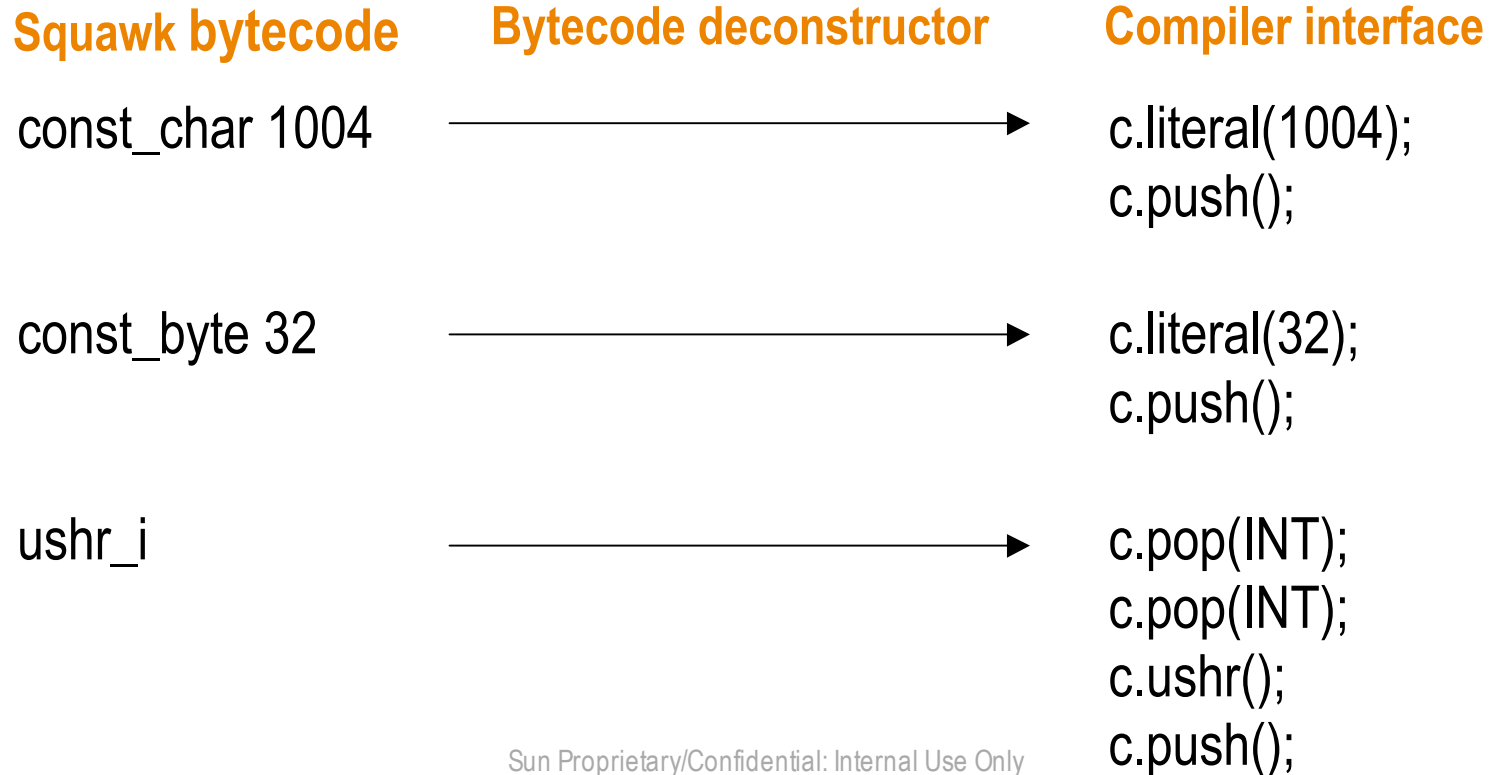
Native code

The Squawk Architecture



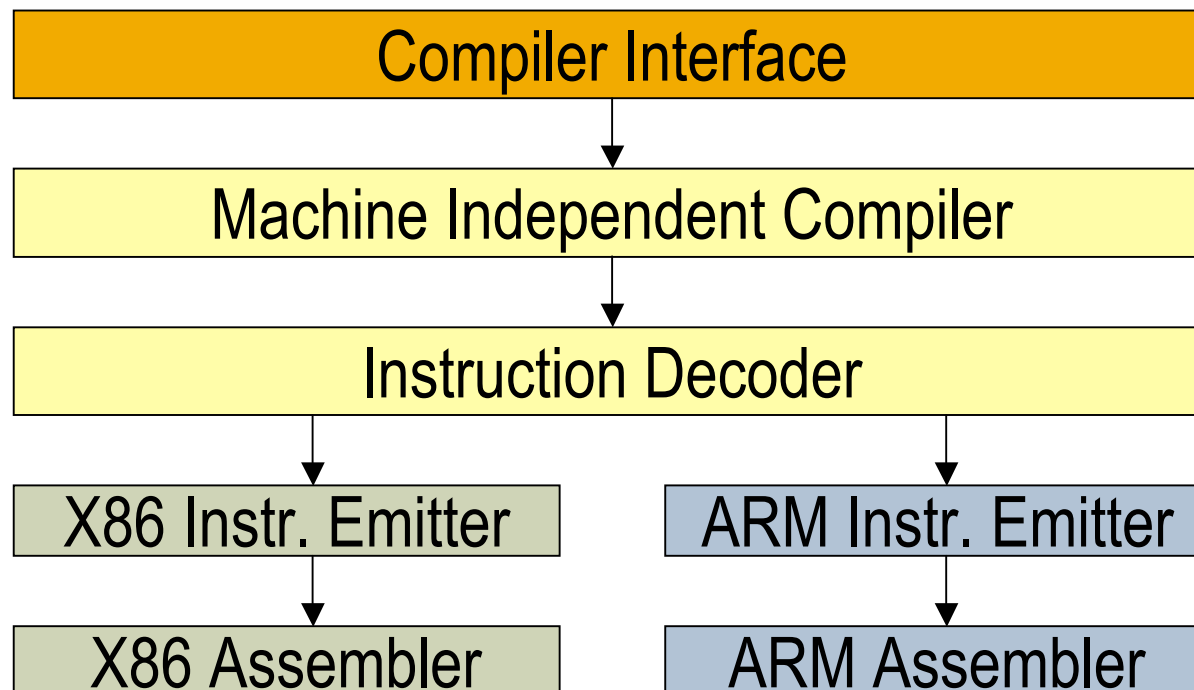
Compiler Interface

- Stack-based, machine-independent instruction interface
- Allows multiple compiler implementations for Squawk
- Two types of compilers will be used by Squawk - simple and optimizing (collaborative research)



Simple Compiler

- Light-weight, single pass, non-optimizing compiler
- Designed for fast code compilation
- Currently supported back-ends - x86, ARM
- Will be used for jitting classes loaded at run time



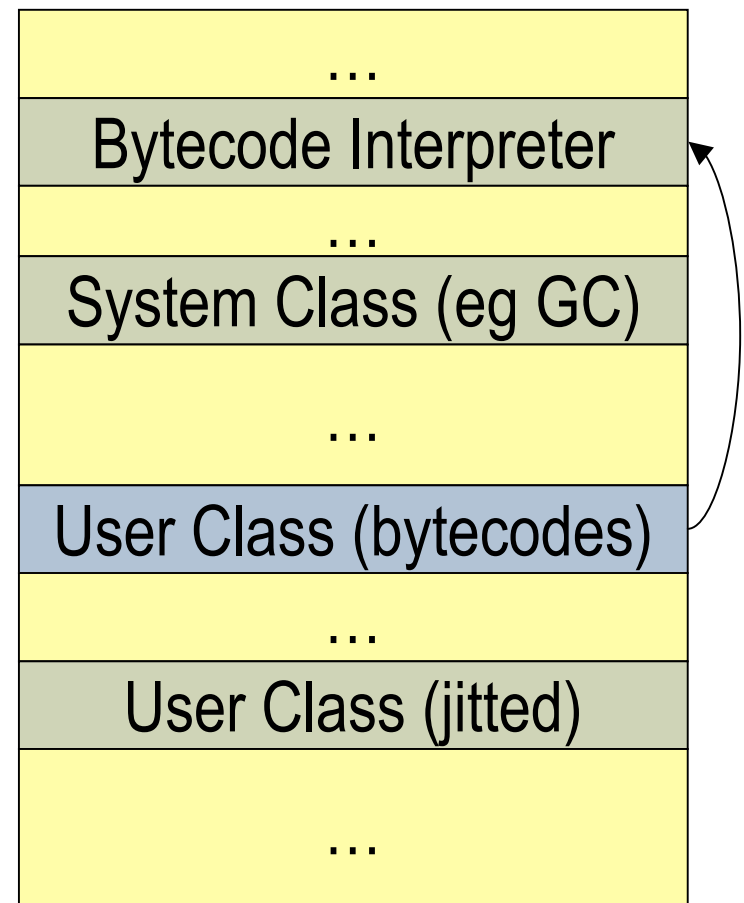
Optimizing Compiler

- Collaborative research with Usyd - Dr Bernhard Scholz
- Aim to test a new instruction selection technique - Partitioned Boolean Quadratic Programming as applied to general purpose machines
- Will be an optimizing compiler:
 - > Constant folding
 - > Dead code elimination
 - > Partial Redundancy Elimination
 - > Etc
- Will be used for:
 - > Ahead-of-time compilation of system code (interpreter, libraries)
 - > Potential jitting of frequently executed code for better performance

Squawk Bytecode Execution

- Currently using a Java interpreter macro-expanded and converted to C
- Future interpreter will be generated through the compiler interface
- Future bytecode execution modes:
 - > Interpreted
 - > Ahead-of-time compiled
 - > Just-in-time compiled
- Mix of bytecodes and compiled code in system memory

System Memory Layout



Misc data

Native code

Bytecodes

Future Work

- Completion of Squawk compiler integration
 - > Testing and debugging the interpreter
 - > Compiling classes ahead-of-time
- Other compiler back-end implementations as needed
- Integration with Scholz's optimizing compiler framework
- Runtime profiling for better jitting
- Running benchmarks with standard JVMs vs Squawk using an optimizing compiler

Summary

- Squawk is a small, execute-in-place CLDC 1.0 compliant Java Virtual Machine
- Written in Java, not C
- Compiler interface for multiple compiler support
- Simple compiler for jitting classes at run time
- Optimizing compiler for ahead-of-time compilation
- Squawk code executed using mix of interpretation, ahead-of-time and just-in-time compilation
- Future work - compiler integration, benchmarking